GPT-5.6 SOL HIGH

Architecture, Compliance & Efficiency Prompt Guide

PHASE 2

PROMPT ARCHITECTURE & CONTROL SYSTEM

DIRECT framework • reasoning selection • validation • defensive cybersecurity control plane

Architecture baseline locked on July 26, 2026

Purpose

Define the reusable prompt operating system for the full guide. Phase 2 establishes how objectives, evidence, roles, boundaries, reasoning configuration, tools, validation, approvals, and output contracts must be expressed before domain prompt packs are authored.

> **INDEPENDENT PUBLICATION BOUNDARY**
> The title is practical shorthand. “High” is a ChatGPT reasoning selection, not a separate base model identity. API reasoning effort and Pro mode are configured independently. This guide does not imply OpenAI endorsement.

Deliverables

- Phase 2 architecture standard and decision matrices
- Copy-ready master prompt templates
- Validation and evaluation controls
- Authorized defensive cybersecurity control system
- Phase 3 production entry criteria

PHASE STATUS

01 Executive control

PHASE 2 STATUS: LOCKED

The prompt architecture is sufficiently defined to govern later leadership, design, code, automation, and defensive cybersecurity prompt packs. Phase 3 may build domain libraries only inside these controls.

Approved architectural decisions

- Use outcome-focused prompts: goal, relevant context, hard constraints, evidence requirements, success criteria, and exact output contract [S01].
- Treat ChatGPT Medium, High, Extra High, and Pro as product-surface selections; treat API model, effort, mode, and verbosity as separate controls [S01][S03].
- Use one accountable role plus explicit specialist lenses instead of title stacking.
- Require evidence and validation proportional to risk; polished prose is never proof of completion.
- Allow safe in-scope inspection and local validation; require confirmation for external writes, destructive actions, purchases, or material scope expansion [S01][S08].
- Keep cybersecurity prompts authorization-first and defensive-outcome-first. Trusted Access does not remove policy, monitoring, authorization, or access controls [S05][S06].
- Use Skills for reusable workflow instructions, plugins as workflow packaging, and apps for connected data/actions subject to existing permissions [S04][S07][S08].

Phase 2 deliverable set

Control	Locked output	Use in later phases
DIRECT	Six-field prompt architecture	Mandatory core of every production prompt
Three tiers	Executive Compact; Production High; Audit-Grade	Right-size prompt length and verification
Reasoning matrix	ChatGPT and API selection logic	Avoid defaulting to maximum effort
Validation gates	Evidence, constraints, tests, release decision	Prevent “looks finished” failures
Cyber control plane	Authorization, scope, evidence, analysis, validation, release	Govern the detailed cyber section
Prompt card schema	Standard metadata and copy block	Make prompt packs consistent and searchable

Sources: S01, S03, S04, S05, S06, S07, S08

CONTROL SURFACE

02 Document map

Phase 2 is organized as a control manual, not a prompt catalog. Each section establishes a reusable design rule that later prompt cards must inherit.

Band	Sections	Production result
A. Model control	Product-surface map; reasoning matrix; selection method	Correct model/effort/mode choices
B. Prompt control	DIRECT; three depth tiers; role architecture	Consistent request construction
C. Context control	Token architecture; evidence; ambiguity; approvals	Efficient, grounded execution
D. Capability control	Skills, plugins, apps, web, files, code, tools	Correct routing and permission boundaries
E. Quality control	Output contracts; validation; evaluation harness	Measurable completion
F. Cyber control	Authorization; risk routing; finding schema; remediation gates	Safe, useful defensive workflows
G. Publishing control	Anti-patterns; prompt card standard; Phase 3 gates	Consistent guide production

READING ORDER

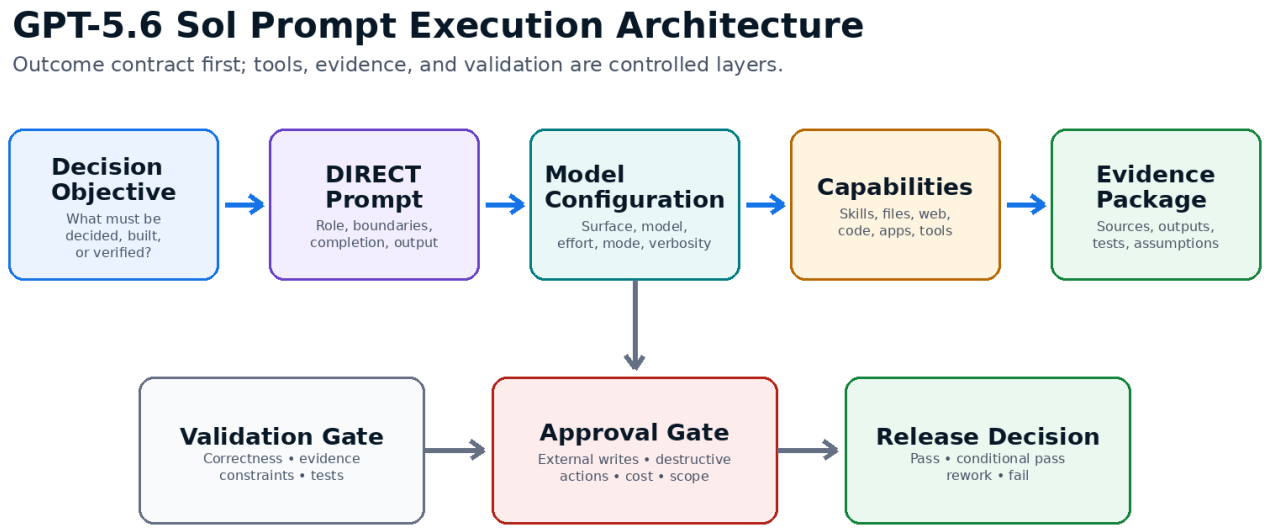
Leaders can begin with Sections 03, 04, 08, 15, and 23. Prompt authors should read the full control system before producing reusable prompt cards. Cybersecurity authors must also use Sections 20–23.

Terminology used throughout

Term	Meaning in this guide
Surface	ChatGPT, Codex, API, or another supported OpenAI product surface.
Configuration	Model selection plus reasoning effort/mode, verbosity, and enabled capabilities.
Capability	A model or product function such as web, files, code execution, Skills, apps, or tool calling.
Prompt	The task-level outcome contract supplied by the user or application.
Validation	Evidence-based checks proving the result meets its completion criteria.
Release	The decision to accept, conditionally accept, rework, or reject the output.

SYSTEM MODEL

03 Architecture at a glance



> **CONTROL PRINCIPLE**
> Do not ask the model to “use everything.” Authorize only the capabilities required by the outcome, provide the evidence they need, and define how their work will be checked.

Execution sequence

1. Define the decision or artifact in observable terms.
2. Supply only evidence and context that can change the result.
3. Choose the surface and reasoning configuration based on measured difficulty.
4. Authorize specific read, analysis, local-change, and validation actions.
5. Require an evidence package, not unsupported confidence.
6. Apply validation and approval gates before external or irreversible action.

Sources: S01, S04, S08

MODEL CONTROL

04 Product-surface and configuration map

The same phrase can mean different controls in ChatGPT and the API. Prompt authors must label the target surface whenever configuration guidance is included.

Surface	User-visible or API control	Verified behavior	Guide rule
ChatGPT	Medium	GPT-5.6 Sol reasoning option	Balanced complex work; availability varies by plan/workspace.
ChatGPT	High	Extended reasoning with GPT-5.6 Sol	Default target for this guide’s complex leadership prompts.
ChatGPT	Extra High	Highest standard reasoning selection with GPT-5.6 Sol	Use when measured quality gains justify additional time.
ChatGPT	Pro	GPT-5.6 Sol Pro product option	Difficult tasks and longer-running workflows; plan/workspace dependent.
API	gpt-5.6-sol or gpt-5.6 alias	Flagship GPT-5.6 model	Choose model separately from effort and mode.
API	reasoning.effort	none, low, medium, high, xhigh, max	Benchmark; never assume max is automatically optimal.
API	reasoning.mode	standard/default or pro execution mode	Pro is independent from model slug and effort.
API	text.verbosity	low, medium, high	Controls default answer detail; task prompt still defines required content.

NAMING RULE

“GPT-5.6 Sol High” is allowed as a practical guide title. Technical passages must state “GPT-5.6 Sol with High reasoning” for ChatGPT, or “gpt-5.6-sol with reasoning.effort: high” for API examples.

Sources: S01, S02, S03

MODEL CONTROL

05 Reasoning-selection matrix

Workload pattern	ChatGPT starting point	API starting point	Escalate when
Routine rewrite, extraction, classification	Instant or Medium	none / low	Errors remain after a clear schema or evidence is supplied.
Standard professional analysis	Medium	medium	Cross-domain synthesis or failure-mode analysis is weak.
Architecture, code review, difficult design	High	high	Representative evaluation shows material omissions.
Quality-critical review with objective checks	High or Extra High	xhigh; compare pro	A marginal reliability gain materially affects the decision.
Hardest quality-first workload	Extra High or Pro	compare xhigh vs max; compare standard vs pro	Only after benchmarking quality, latency, and cost.
High-volume repeatable workflow	Lower level plus fixed schema	low / medium; Terra or Luna may fit	Exceptions or review failures exceed tolerance.

Selection rule

```
1. Establish the lowest plausible baseline.
2. Test representative tasks with the same configuration and one level lower.
3. Score correctness, completeness, evidence, constraints, latency, tokens, and cost.
4. Escalate only where the measured quality gain matters.
5. Preserve the prompt; do not add “think harder” instructions to simulate a configuration.
```

OFFICIAL EFFICIENCY GUIDANCE

OpenAI recommends preserving the current effort as a baseline and testing one level lower because GPT-5.6 can often maintain or improve quality with fewer tokens. Max and Pro should be evaluated, not assumed.

Sources: S01, S03

MODEL CONTROL

06 Configuration decision method

Use the following six questions before prescribing High, Extra High, max, or Pro. A high-stakes label alone does not prove that more model work improves the outcome.

Question	Low-cost answer	Escalation signal
Is the output objectively testable?	Use schema and deterministic checks first.	More reasoning improves verified success rate.
Does the task require multi-domain synthesis?	Medium may be sufficient.	High materially reduces missed dependencies.
Is fresh judgment needed between tool calls?	Use direct or programmatic tools for bounded work.	Agentic reasoning is required after each result.
Is latency operationally important?	Keep a lower-effort baseline.	Quality gain outweighs response-time impact.
Will the result trigger an external action?	Add approval gates before escalating effort.	The decision value justifies deeper verification.
Can the task be cleanly decomposed?	Use a single thread when tightly coupled.	Parallel independent workstreams justify multi-agent beta.

Configuration record

```
Surface: [ChatGPT / API / Codex]
Model or product option: [High / Extra High / Pro / gpt-5.6-sol]
Reasoning effort: [API only: none | low | medium | high | xhigh | max]
Reasoning mode: [API only: standard | pro]
Verbosity: [API only: low | medium | high]
Enabled capabilities: [only those required]
Evaluation baseline: [task set and acceptance threshold]
Escalation rule: [measured condition]
```

Sources: S01

PROMPT CONTROL

07 DIRECT prompt architecture

DIRECT is the mandatory six-field control plane for production prompts. It replaces vague persona inflation with observable responsibilities, evidence, boundaries, and completion conditions.

Field	Control question	Minimum required content
D — Decision objective	What must be decided, built, diagnosed, or verified?	One observable outcome; target audience or operator.
I — Inputs and evidence	What may be used, and what is authoritative?	Files, data, code, sources, dates, provenance, missing evidence.
R — Role and responsibility	Who is accountable, and which lenses apply?	One accountable role; specialist lenses; conflict priority.
E — Execution boundaries	What actions, assumptions, or changes are forbidden?	Scope, approvals, safe local actions, external-action gates.
C — Completion criteria	What proves the task is finished?	Tests, evidence, quality thresholds, residual issues.
T — Target output contract	What exact form must be returned?	Sections, schema, file type, length, ordering, citations.

MINIMUM VIABLE PROMPT

A short prompt can still be production-grade when all six controls are either explicit or safely inherited from a stable Skill, project brief, or application layer. Do not repeat inherited rules without reason.

```
D: Produce [observable outcome] for [audience].
I: Use [authoritative inputs]. State material evidence gaps.
R: Act as the accountable [role]; apply [specialist lenses].
E: Stay within [scope]. Do not [prohibited actions]. Confirm before [gated actions].
C: Complete only when [tests/evidence/acceptance conditions].
T: Return [exact structure, schema, file, length, and ordering].
```

Sources: S01

DIRECT CONTROL

08 D and I: objective, inputs, and evidence

D — Decision objective

Write the objective as a verifiable state change, artifact, or decision. “Help with the architecture” is not a completion target. “Produce a target-state architecture and ADR set that resolves the listed integration and tenancy decisions” is.

Weak form	Controlled form
Review this code.	Identify release-blocking security and reliability defects in the supplied repository; provide file/line evidence, patch guidance, and regression tests.
Make this design better.	Redesign the supplied onboarding screen to reduce cognitive load while preserving required fields, brand constraints, and accessibility; return annotated decisions and implementation-ready specifications.
Plan our integration.	Produce a target-state integration architecture with trust boundaries, error handling, idempotency, observability, migration stages, and decision records.

I — Inputs and evidence

- Name authoritative sources and their order of precedence.
- Separate supplied facts from assumptions, external research, and model inference.
- Provide dates for volatile facts and require current verification when necessary.
- State what evidence is unavailable; do not force a confident conclusion from incomplete inputs.
- For files and code, require precise references such as section, page, path, symbol, or line where available.
- Treat connected data as permission-scoped evidence, not as universal access.

> **EVIDENCE PRIORITY**
> Primary source or supplied artifact > current official documentation > authoritative standard > reputable secondary source > clearly labeled inference.

DIRECT CONTROL

09 R and E: accountability and execution boundaries

R — Role and responsibility

Role prompting is useful when it specifies decision authority, review lens, risk tolerance, and professional standard. A title alone contributes little.

```
Accountable role: Director of Solution Architecture
Specialist lenses: application security, data architecture, SRE, UX systems
Decision priority: safety and data integrity > correctness > maintainability > speed
Conflict rule: identify unresolved tradeoffs; do not silently average incompatible requirements
Excluded authority: legal approval, production deployment, vendor purchasing
```

E — Execution boundaries

Boundary class	Default rule
Read / inspect	Allowed when relevant to answering, reviewing, diagnosing, or planning.
Local build / fix	Allowed when explicitly requested; run relevant non-destructive validation.
External write	Require confirmation unless the user explicitly authorized the exact action and product permissions permit it.
Destructive / irreversible	Require confirmation and a rollback or recovery plan.
Cost / purchase	Require confirmation before incurring cost or committing spend.
Scope expansion	Stop or request approval when the change is material to risk, time, data, or systems.
Sensitive data	Use only what is needed; preserve access and confidentiality boundaries.

BOUNDARY PRECISION

State each approval rule once. Repeating “ask first” throughout a prompt can cause unnecessary pauses for safe local work.

Sources: S01, S08

DIRECT CONTROL

10 C and T: completion and output contract

C — Completion criteria

Completion must be evidence-based. Replace “make it perfect” with observable acceptance conditions.

Artifact type	Completion examples
Code change	Target tests pass; lint/type checks pass; no unintended diff; rollback documented; known failures disclosed.
Architecture	Requirements mapped; trust boundaries shown; failure modes addressed; decisions and unresolved risks recorded.
Research	Claims cited; freshness checked; disagreement represented; inference labeled; unanswered questions listed.
Design	Required states covered; hierarchy and accessibility reviewed; responsive behavior specified; visual QA completed.
Cyber review	Authorization documented; findings evidenced; false positives considered; remediation and regression tests supplied.

T — Target output contract

```
Return exactly:
1. Executive decision — 5 sentences maximum.
2. Findings — severity order; maximum 10; evidence required.
3. Recommended action — owner, sequence, dependencies, rollback.
4. Validation performed — tests, checks, and results.
5. Open risks — assumptions, missing evidence, residual risk.
Format: Markdown plus [requested artifact file].
Do not include generic praise, hidden-reasoning requests, or unsupported certainty.
```

OUTPUT DISCIPLINE

The output contract controls what the user receives; it should not micromanage private reasoning. Ask for conclusions, evidence, assumptions, tests, and unresolved risks.

PROMPT CONTROL

11 Three prompt-depth tiers

Tier	Best use	Typical structure	Configuration guidance
1 — Executive Compact	Experienced operator; stable context; low ambiguity	Six compact DIRECT lines	Medium or High; evaluate lower effort for repeat use.
2 — Production High	Complex professional task; files/tools; measurable output	Full DIRECT with priorities and validation	High starting point; Extra High/xhigh only if benchmarked.
3 — Audit-Grade	Quality-critical decision; release gate; regulated or security-sensitive work	DIRECT plus provenance, assumptions, risks, tests, rollback, traceability	Pro-eligible; compare standard/pro and xhigh/max on representative tasks.

Tier selection

- Use the shortest tier that preserves all material controls.
- Escalate prompt depth when ambiguity, evidence complexity, action risk, or validation burden increases.
- Do not confuse a long prompt with a difficult task; configuration and prompt length are independent.
- Stable instructions belong in a Skill or project brief; task prompts should reference them rather than repeat them.
- Audit-Grade is a prompt architecture. It does not automatically require ChatGPT Pro or API pro mode.

> **TOKEN-CONSERVATIVE RULE**
> A production prompt should be concise enough to audit. Repetition, decorative role language, generic warnings, and unused background consume context without improving control.

PROMPT TEMPLATE

12 Tier 1 — Executive Compact

Use when the operator already has the necessary context and the task can be fully controlled in a compact form.

```
Act as the accountable [role].

D — Objective: [observable decision or artifact].
I — Evidence: Use [inputs]. Treat [source] as authoritative. State material gaps.
R — Lenses: Apply [specialist lenses] in this priority order: [order].
E — Boundaries: Stay within [scope]. Do not [prohibited actions]. Confirm before [gated actions].
C — Complete when: [tests, evidence, acceptance threshold].
T — Return: [exact sections, schema, length, order, and file type].
```

Example: architecture decision

```
Act as the accountable solution architect.

D — Objective: Decide whether the supplied webhook workflow should remain synchronous or move to an
event queue.
I — Evidence: Use the architecture diagram, failure logs, throughput data, and current provider
documentation. State missing measurements.
R — Lenses: Data integrity > reliability > operability > cost.
E — Boundaries: Do not change production or invent provider guarantees.
C — Complete when: The recommendation addresses idempotency, retries, ordering, dead letters,
observability, migration, and rollback.
T — Return: One-page decision memo plus an ADR in Markdown.
```

> **COMPACT DOES NOT MEAN VAGUE**
> Every line does control work. Remove any line only when the rule is safely inherited from a stable context layer.

PROMPT TEMPLATE

13 Tier 2 — Production High

```
ROLE
You are the accountable [role] for this work. Apply [specialist lenses]. When priorities conflict, use:
[priority order].

MISSION
Produce [observable artifact or decision] for [audience/operator].

CONTEXT AND AUTHORIZED EVIDENCE
Use [files, code, data, official sources, connected systems]. Treat [sources] as authoritative.
Distinguish evidence, assumptions, and inference.

EXECUTION BOUNDARIES
Stay within [scope]. You may [safe reads/local actions/tests]. Do not [prohibited actions]. Require
confirmation before [external, destructive, costly, or scope-expanding actions].

DECISION REQUIREMENTS
Address: [required dimensions]. Rank findings by [method].

VALIDATION
Verify with [tests, render, citations, calculations, diff, source check]. Report failures and residual
uncertainty.

AMBIGUITY POLICY
Ask only when a missing fact would materially change the result or authorization. Otherwise state the
assumption and continue.

OUTPUT CONTRACT
Return [exact structure, schema, length, order, artifact files].

COMPLETION RULE
Do not declare completion until [objective criteria].
```

HIGH REASONING USE

Select High because the workload benefits from extended reasoning, not because the prompt contains the word “High.” Configuration is chosen in the product or API.

Sources: S01, S03

PROMPT TEMPLATE

14 Tier 3 — Audit-Grade, Pro-eligible

Use for quality-critical architecture, release, security, due-diligence, compliance, and high-value engineering decisions. Add only controls that can be verified.

```
ACCOUNTABILITY
Accountable role: [role]. Review lenses: [lenses]. Decision authority and exclusions: [authority].

DECISION OBJECTIVE
[Observable outcome, audience, decision deadline, and consequence of error.]

EVIDENCE REGISTER
Authorized sources: [list]. Authority order: [order]. Freshness requirement: [date/rule]. Missing evidence:
[known gaps].

SCOPE AND APPROVALS
In scope: [assets/actions]. Out of scope: [exclusions]. Safe local actions: [list]. Approval required for:
[external/destructive/cost/scope].

METHOD AND TRACEABILITY
Map every material conclusion to evidence. Maintain an assumption register, risk register, and decision log.
Identify conflicting evidence.

FAILURE-MODE REVIEW
Evaluate [technical, operational, security, compliance, data, human, and rollback failure modes].

VALIDATION PLAN
Run [tests/checks]. Define pass/fail thresholds. Record actual results, exceptions, and untested areas.

RELEASE CONTRACT
Return: executive decision; evidence map; prioritized findings; recommended plan; validation record; residual
risks; rollback; unresolved decisions.

COMPLETION
Status must be PASS, CONDITIONAL PASS, REWORK, or FAIL. Never report PASS with an unmet mandatory criterion.
```

PRO-ELIGIBLE, NOT PRO-REQUIRED

Benchmark the same prompt in standard and Pro configurations. Use Pro only when its measured reliability gain justifies additional latency and token use.

Sources: S01

PROMPT CONTROL

15 Role architecture

The role field is an accountability model. It should tell the model whose decision it is producing, which disciplines must be applied, and how conflicts are resolved.

Component	Required definition	Failure prevented
Accountable role	Single owner of the decision or artifact	Competing voices and unfocused output.
Specialist lenses	Named disciplines that review the work	Missing cross-functional constraints.
Priority order	Explicit order for tradeoffs	Silent averaging of incompatible goals.
Decision rights	What the model may recommend, edit, or execute	Unauthorized action or overreach.
Exclusions	What still requires human/legal/production approval	False authority and compliance claims.
Audience	Who will consume or operate the result	Wrong level of detail or unusable format.

Recommended pattern

```
Accountable role: Director of [discipline]
Specialist lenses: [2–5 relevant disciplines]
Priority order: [non-negotiable safety/correctness priorities first]
Decision rights: [recommend / edit local artifact / test / no external write]
Excluded authority: [legal approval / production deployment / purchase / policy waiver]
Audience: [executive / engineer / operator / customer / auditor]
```

AVOID TITLE INFLATION

“Director, architect, engineer, designer, compliance officer, and hacker” is not a role contract. Use one accountable role and assign the others as review lenses with explicit priorities.

PROMPT TEMPLATE

16 Role-stack composer

Workload	Accountable role	Useful lenses	Default priority order
Solution architecture	Director of Solution Architecture	AppSec, data, SRE, integration, FinOps	Safety/data integrity > correctness > reliability > maintainability > cost
Code remediation	Engineering lead	Secure coding, testing, performance, operability	Correctness > security > regression safety > maintainability > speed
Product design	Design director	UX research, accessibility, frontend, conversion	User safety > task success > accessibility > clarity > aesthetics
Cyber review	Security architecture lead	AppSec, cloud/IAM, detection, incident response	Authorization > containment > evidence > remediation > speed
Executive decision	CTO or CISO delegate	Architecture, finance, operations, compliance	Risk-adjusted value > feasibility > time > cost

Conflict-handling clause

```
When specialist lenses disagree:
1. Identify the conflict explicitly.
2. Apply the stated priority order.
3. Preserve non-negotiable constraints.
4. Quantify or characterize the tradeoff.
5. Escalate only when the choice exceeds the role’s authority or materially changes scope.
```

Role-stack prompt fragment

```
You are the accountable [role]. Consult the work through [lens 1], [lens 2], and [lens 3]. Do not simulate separate personalities or produce redundant opinions. Synthesize one decision. Resolve conflicts using [priority order], and flag any decision that remains outside your stated authority.
```

EFFICIENCY CONTROL

17 Context and token architecture

Keep stable controls separate from changing task evidence. This improves auditability, reduces repetition, and makes prompt caching and reusable Skills more effective.

Layer	Contents	Change rate	Efficiency rule
1. System / workspace	Safety, product, organization, and access controls	Rare	Do not restate or attempt to override.
2. Skill / reusable workflow	Stable procedure, standards, templates, validation	Low	Load only when relevant; keep core instructions compact.
3. Project brief	Architecture, brand, repository, definitions, long-lived decisions	Medium	Reference a maintained source of truth.
4. Task evidence	Files, logs, code, data, current research	High	Include only evidence that can change the result.
5. Output contract	Schema, length, ordering, artifact requirements	High	State once and make it testable.
6. Validation state	Tests, render results, failed checks, open risks	High	Carry forward only what remains relevant.

Token waste to remove

- Repeated persona titles and motivational language.
- Duplicated constraints expressed in multiple sections.
- Generic background that does not affect the decision.
- Full logs or files when a bounded relevant range is sufficient.
- Requesting multiple drafts when one answer plus a validation gate is enough.
- Long explanations of obvious formatting rules instead of an explicit schema.

> **PROGRESSIVE LOADING**
> Skills are designed to carry reusable instructions and supporting resources, allowing detailed references or code to be used only when the task requires them.

Sources: S01, S04

EFFICIENCY CONTROL

18 API efficiency controls

These controls are API-specific. They must not be presented as ordinary ChatGPT prompt syntax.

Control	Use when	Guardrail
reasoning.effort	The task benefits from measured additional reasoning	Compare the same level and one lower; reserve max for hardest workloads.
reasoning.mode: pro	A marginal reliability gain materially affects a difficult outcome	Benchmark against standard mode; tokens and latency increase.
text.verbosity	A stable default detail level is useful across requests	Prompt still states mandatory content and structure.
reasoning.context: all_turns	Goals, assumptions, and priorities remain stable across turns	Do not carry obsolete reasoning into a changed task.
reasoning.context: current_turn	Earlier reasoning is no longer relevant	Preserve needed evidence explicitly.
Prompt caching	A reusable prefix is stable and repeated	Track cache writes/reads; avoid unnecessary writes.
Programmatic Tool Calling	Bounded tool-heavy processing needs little fresh judgment between calls	Measure final-answer completeness, not just fewer calls.
Multi-agent beta	Independent workstreams can run in parallel and be synthesized	Avoid for tightly coupled decisions or shared mutable state.

EFFICIENCY METRIC

Fewer tokens or tool calls are improvements only when the final answer still satisfies correctness, evidence, and completion requirements.

Sources: S01

GROUNDING CONTROL

19 Evidence and source protocol

Evidence class	Examples	Required treatment
Supplied primary artifact	Code, contract, logs, architecture, spreadsheet, PDF	Inspect directly; cite precise location; do not substitute web summaries.
Current official source	OpenAI docs, vendor docs, laws, standards publisher	Verify freshness; capture date/version; cite claim inline.
Connected internal source	Drive file, email, calendar, CRM, repository	Respect source permissions; distinguish retrieved content from inference.
Deterministic output	Test result, calculation, render, diff, parser output	Record command/method and actual result.
Secondary source	News, analysis, community content	Use for context; do not outrank primary sources without explanation.
Inference	Conclusion derived from multiple facts	Label explicitly; cite supporting evidence; state uncertainty.
Assumption	Unverified condition needed to proceed	Record it; explain impact; validate if material.

Evidence clause

```
For every material claim, identify whether it is: supplied fact, current official source,
deterministic test result, connected internal evidence, inference, or assumption. Cite the
smallest useful source location. When evidence conflicts, present the conflict and explain the
decision rule. Do not convert missing evidence into confidence.
```

CURRENT-INFORMATION RULE

Model knowledge is not a substitute for current verification when facts may have changed. Use current official sources for product behavior, policies, software documentation, and standards.

GROUNDING CONTROL

20 Ambiguity and assumption protocol

The goal is neither “never ask questions” nor “ask about everything.” The model should pause only when the missing fact materially affects correctness, authorization, safety, or scope.

Condition	Required behavior
Authorization or target ownership is unclear	Ask before advanced cyber analysis or any external action.
Two interpretations produce materially different deliverables	Ask one focused question or present the alternatives and request selection.
A non-material detail is missing	State a reasonable assumption and continue.
Current fact may have changed	Verify from an authoritative current source.
Evidence is incomplete but a bounded analysis is still useful	Proceed with limitations, confidence, and missing-evidence list.
Required source or file is unavailable	Do not invent its contents; identify the gap and the next retrieval step.
The task exceeds authorized scope	Stop at the boundary and request explicit expansion.

Assumption register

```
A-01 | Assumption | Why needed | Evidence for/against | Impact if false | Validation owner |
Status
A-02 | ...
```

QUESTION ECONOMY

Ask the smallest question that resolves the material ambiguity. Do not return a questionnaire when one missing field controls the decision.

ACTION CONTROL

21 Autonomy and approval boundaries

Request type	Default autonomy	Approval boundary
Answer / explain	Inspect relevant materials and answer	No implementation unless requested.
Review / diagnose	Read, analyze, test safely, and report	Do not mutate external systems.
Plan / design	Produce decisions, artifacts, and validation plan	No execution of the plan without authorization.
Build / fix local artifact	Make requested in-scope changes; run non-destructive tests	Confirm before external deployment or destructive migration.
Connected-app read	Read when enabled and relevant, subject to permissions	Respect app and source-system access.
Connected-app write	Prepare the action and use product approval controls	Confirm for meaningful external effects unless explicitly pre-authorized.
Purchase / spend / subscription	Research and recommend	Confirm before commitment.
Scope expansion	Explain why it is needed	Confirm before material expansion.

Compact approval policy

```
For answer, review, diagnose, or plan requests: inspect relevant materials and report the result;
do not implement changes unless requested.

For change, build, or fix requests: make in-scope local changes and run relevant non-destructive
validation.

Require confirmation for external writes, destructive actions, purchases, or material scope
expansion.
```

Sources: S01, S08

CAPABILITY CONTROL

22 Skills, plugins, apps, and tool routing

Capability	Use for	Do not assume
Skill	Reusable instructions, examples, scripts, templates, organization conventions	That the Skill is installed, enabled, trusted, or available on every surface.
Plugin	Packaged workflow that may contain Skills, apps, and app templates	That installing it grants access to underlying systems.
App	Connected data, tools, UI, or actions	That app permissions override source-system permissions.
Files	Supplied or retrieved documents, code, and datasets	That a filename or snippet contains the full relevant content.
Web research	Current public facts and official documentation	That search snippets are sufficient evidence without opening the source.
Code runtime	Deterministic computation, parsing, tests, rendering, artifact generation	That execution results are correct without validation and inspection.
Human approval	External, destructive, costly, sensitive, or scope-expanding decisions	That a model recommendation transfers accountability.

Routing instruction

```
Select the minimum capability set that can complete and validate the task. Prefer direct evidence over summaries, deterministic checks over verbal confidence, and existing permissions over inferred access. If a required capability is unavailable, state the limitation and do not claim it was used.
```

PERMISSION INHERITANCE

Plugins and apps remain subject to plan, workspace, role, region, connection status, action controls, and the user's permissions in the source system.

Sources: S04, S07, S08, S09

QUALITY CONTROL

23 Output-contract patterns

Need	Contract pattern
Executive decision	Decision; rationale; top risks; immediate actions; owner; deadline; confidence.
Architecture	Context; requirements; diagram; data/trust boundaries; decisions; failure modes; migration; observability; risks.
Code review	Finding ID; severity; file/line; evidence; preconditions; impact; patch; test; confidence.
Implementation	Plan; files changed; change summary; validation commands/results; rollback; unresolved issues.
Research	Answer; evidence; source date; disagreement; inference; limitations; next verification.
Design	User goal; hierarchy; states; responsive behavior; accessibility; visual specification; QA findings.
Cyber finding	Authorization; asset; evidence; impact; likelihood; confidence; remediation; regression test; residual risk.

Structured-output fragment

```
Return a JSON-compatible object with these fields only:
{
  "decision": "...",
  "findings": [{"id":"F-01","severity":"...","evidence":"...","action":"..."}],
  "validation": [{"check":"...","result":"pass|fail|not_run","evidence":"..."}],
  "assumptions": ["..."],
  "residual_risks": ["..."],
  "status": "PASS|CONDITIONAL_PASS|REWORK|FAIL"
}
```

SCHEMA CAUTION

Use a strict schema when downstream automation needs determinism. Use prose when judgment, nuance, or stakeholder communication is the primary output.

QUALITY CONTROL

24 Validation-gate system

Gate	Question	Evidence required
G1 — Scope	Was only the authorized task performed?	Scope map; changed files/actions; exclusions.
G2 — Correctness	Does the result solve the stated objective?	Tests, calculations, source checks, acceptance criteria.
G3 — Evidence	Are material conclusions traceable?	Citations, file/line references, logs, commands, screenshots or renders.
G4 — Constraints	Were security, compliance, brand, format, and action boundaries followed?	Constraint checklist and exceptions.
G5 — Failure modes	Were meaningful failure, abuse, rollback, and edge cases reviewed?	Risk register and mitigation mapping.
G6 — Presentation	Is the artifact usable, legible, and internally consistent?	Render inspection, schema validation, link check.
G7 — Release	Can the output be accepted at the required risk level?	PASS / CONDITIONAL PASS / REWORK / FAIL with rationale.

Release statuses

Status	Rule
PASS	All mandatory criteria met; no unresolved release blocker.
CONDITIONAL PASS	Usable only with explicit conditions, owner, and deadline.
REWORK	Material defects remain, but the path to acceptance is defined.
FAIL	Objective not met, evidence insufficient, boundary violated, or risk unacceptable.

NO CEREMONIAL VALIDATION

“Reviewed carefully” is not a result. Name the check, method, actual output, and unresolved coverage.

QUALITY CONTROL

25 Evaluation and benchmark harness

Dimension	Weight	Scoring question
Correctness	25	Is the core decision or artifact technically right?
Evidence	15	Are material claims traceable to sufficient evidence?
Completeness	15	Are all mandatory requirements and failure modes addressed?
Constraint adherence	15	Were scope, security, format, and approval boundaries followed?
Risk handling	10	Are uncertainty, edge cases, and residual risks treated correctly?
Actionability	10	Can the intended operator use the result without reconstruction?
Efficiency	5	Does the configuration avoid unnecessary tokens, calls, or latency?
Presentation	5	Is the output clear, legible, consistent, and correctly formatted?

Score	Decision
90–100	Release candidate, subject to mandatory gate results.
80–89	Conditional pass or targeted rework.
70–79	Rework; configuration or prompt likely needs adjustment.
Below 70	Fail; do not release.

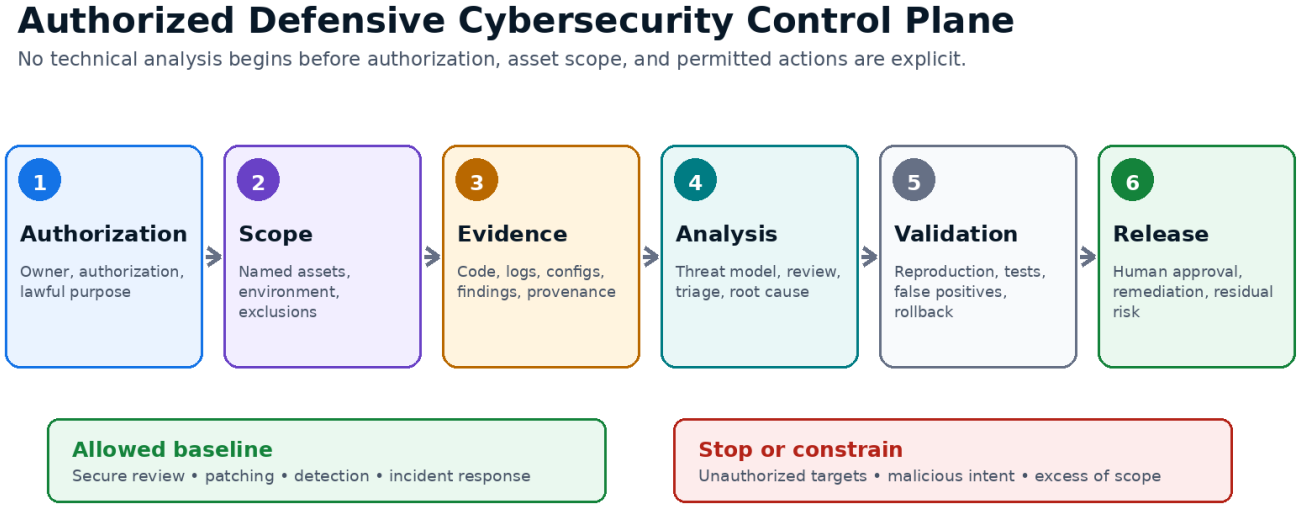
Benchmark method

1. Build a representative task set with known acceptance criteria and difficult edge cases.
2. Run the same prompt and evidence across candidate configurations.
3. Blind-score outputs where practical; record tokens, latency, tool calls, and cost.
4. Select the lowest-cost configuration that reliably clears the release threshold.
5. Re-test after model, prompt, Skill, tool, or source changes.

Sources: S01

CYBER CONTROL

26 Defensive cybersecurity control plane



MANDATORY BOUNDARY

No cyber prompt in the final guide may omit authorization, named scope, permitted actions, prohibited actions, evidence, and completion conditions when the work could materially affect a system or facilitate dual-use activity.

Baseline allowed defensive outcomes

- Secure code review, threat modeling, patch development, debugging, and defensive testing.
- Human-led vulnerability identification and validation in authorized systems.
- Corporate security automation, detection engineering, incident-response support, and hardening.
- Security education and analysis that does not facilitate malicious or unauthorized activity.

Sources: S05, S06, S10, S11

CYBER TEMPLATE

27 Cyber authorization header

```
AUTHORIZED DEFENSIVE CONTEXT

System owner or authorizing organization: [name]
Operator and authority basis: [owner / employee / contract / written authorization]
Authorized target: [repository, application, account, network, lab, or dataset]
Environment: [production / staging / local lab / CTF / isolated test]
Defensive objective: [review, validate, patch, detect, contain, recover, document]
Permitted actions: [read, inspect, reproduce in lab, test, patch, scan bounded scope]
Prohibited actions: [third-party targeting, persistence, exfiltration, destructive change, uncontrolled
deployment]
Data handling limits: [sensitivity, retention, redaction, export limits]
Stop conditions: [unexpected target, production instability, sensitive discovery, scope uncertainty]
Evidence available: [code, logs, configurations, scan output, architecture, prior findings]
Completion condition: [finding validated, patch tested, detection verified, risk accepted]
Human approval required for: [production change, disclosure, external communication, destructive action]
```

Authorization handling rules

- Do not treat the header as a legal authorization document; it records the user-provided scope for execution.
- Ask for clarification when ownership, target, environment, or permitted actions are materially ambiguous.
- Stay inside granted authorization even when a technically adjacent action appears useful.
- Preserve evidence and avoid destructive validation when a safer reproduction is available.
- Trusted Access supports authorized defenders but does not waive usage policies, monitoring, or approved scope.

Sources: S05, S06

CYBER CONTROL

28 Cyber risk routing and boundaries

Route	Examples	Prompt treatment
GREEN — baseline	Threat modeling; secure code review; patching; detection; hardening; incident planning	Proceed with authorization and evidence controls.
AMBER — dual use	Exploit reproduction; malware analysis; credential simulations; chained research; production attack-path testing	Require explicit authorization, bounded targets, safe environment, stop conditions, and defensive outcome. Constrain detail if needed.
RED — excluded	Unauthorized targeting; credential theft; exfiltration; destructive access; malicious persistence; indiscriminate malware	Do not provide operational assistance. Redirect to safe defensive, recovery, or reporting guidance.

Risk-routing questions

1. Is the target owned, operated, or explicitly authorized for testing?
2. Is the objective defensive and bounded, or does it facilitate compromise beyond the stated scope?
3. Can the same defensive result be achieved with a safer lab, code review, simulation, or detection-focused method?
4. Would the requested detail materially increase the ability to harm third parties or evade controls?
5. Are production safety, sensitive data, disclosure, and stop conditions defined?

> **SAFEGUARD INTERACTION**
> Additional automated checks may delay or block some cyber requests. A check does not by itself establish a policy violation. Clarifying legitimate authorization and defensive scope is appropriate; disguising intent or attempting to bypass safeguards is not.

Sources: S05, S06, S10, S11

CYBER CONTROL

29 Cyber finding and evidence schema

Field	Required content
Finding ID and title	Stable identifier and precise defect statement.
Authorization / asset	Authorized target, environment, and affected component.
Severity and priority	Defined method; business context; avoid severity inflation.
Evidence	File/line, request/response, log, configuration, trace, or reproducible observation.
Preconditions / reachability	What must be true for the issue to be reachable or exploitable.
Impact	Confidentiality, integrity, availability, safety, compliance, or business effect.
Likelihood / confidence	Separate likelihood from confidence; explain uncertainty.
False-positive analysis	Alternative explanations and checks performed.
Remediation	Specific code, configuration, control, or process change.
Regression validation	Test case, expected result, negative test, monitoring.
Residual risk	What remains after remediation and who accepts it.
Disclosure / handling	Sensitivity, owner, communication, and approval requirements.

Secure-code-review output fragment

```
For each finding return:
ID | severity | CWE/control mapping if useful | file:line | evidence | reachability | impact |
likelihood | confidence | false-positive checks | recommended patch | regression test | residual
risk

Do not report a vulnerability solely from a pattern match. Trace relevant data/control flow and
state what was or was not reproduced.
```

CYBER CONTROL

30 Cyber validation and remediation gate

Gate	Mandatory evidence
Authorization	Named owner/authority, target, environment, permitted and prohibited actions.
Reproduction	Safe reproduction or sufficient code/configuration proof; no uncontrolled production testing.
Root cause	Specific vulnerable logic, trust failure, configuration, or control gap.
Impact validation	Reachability and realistic impact; not only theoretical severity.
Remediation	Patch or control mapped to root cause, with side effects considered.
Regression	Positive and negative tests; affected variants; monitoring or detection where relevant.
Deployment safety	Sequence, backup, rollback, maintenance impact, and approvals.
Residual risk	Unresolved paths, compensating controls, owner, and acceptance status.

Remediation completion clause

```
Do not mark the finding remediated because code was changed. Mark it remediated only when the root cause is addressed, the intended security property is tested, relevant regressions are checked, deployment/rollback is defined, and residual risk is recorded. If a required validation cannot be run, status is CONDITIONAL PASS or REWORK—not PASS.
```

DEFENDER ADVANTAGE

OpenAI reports that GPT-5.6 is presently stronger at finding and fixing vulnerabilities than at reliably executing autonomous end-to-end attacks against hardened targets. The guide therefore concentrates advanced capability on evidence-based remediation and defensive operations.

Sources: S05

PUBLISHING CONTROL

31 Prompt anti-patterns and corrections

Anti-pattern	Why it fails	Correction
Title stacking	Creates competing personas without decision rights.	One accountable role plus specialist lenses and priority order.
“Think step by step” as a quality strategy	Does not replace configuration, evidence, or validation.	Select reasoning in the product/API; request conclusions, evidence, tests, and risks.
Maximum effort by default	Can waste latency/tokens without measured gain.	Benchmark the same setting and one lower; compare xhigh/max and standard/pro.
“Use all tools”	Expands scope and increases irrelevant work.	Authorize the minimum capability set required.
Repeated warnings	Consumes context and can cause unnecessary pauses.	State each boundary once in the execution section.
Never ask questions	Forces unsafe assumptions in material ambiguity.	Ask only when correctness, authorization, or scope changes.
No output schema	Produces verbose or unusable answers.	Define exact sections, fields, ordering, and length.
No completion test	Encourages premature “done” claims.	Define objective acceptance and validation gates.
Connected app = permission	Misstates access and can expose sensitive assumptions.	Respect app, workspace, and source-system permissions.
Cyber euphemisms to bypass safeguards	Obscures intent and undermines legitimate defense.	State authorization and defensive purpose directly.

Sources: S01, S04, S05, S08

PUBLISHING CONTROL

32 Prompt-card production standard

Card field	Requirement
ID and title	Stable identifier; outcome-oriented name.
Use case	When to use and when not to use.
Surface	ChatGPT, API, Codex, or surface-agnostic.
Recommended configuration	Starting model/reasoning; benchmark/escalation note.
Required inputs	Files, evidence, authorization, variables.
Copy-ready prompt	DIRECT-compliant template with replaceable fields.
Variable map	Definitions, allowed values, examples.
Output contract	Exact sections, schema, artifact, length.
Validation	Checks, tests, evidence, release threshold.
Tool/Skill associations	Only verified or conditional capabilities; no universal-access claim.
Token note	What can be inherited, cached, omitted, or lowered.
Security/compliance note	Approval, privacy, scope, or cyber boundary.
Source note	Official source IDs supporting product-specific claims.

Card status

Status	Meaning
DRAFT	Prompt written but not tested.
TESTED	Run on representative examples with recorded results.
RELEASED	Meets threshold and editorial review.
CONDITIONAL	Requires a named environment, app, Skill, authorization, or plan.
DEPRECATED	Superseded or no longer accurate.

RELEASE CONTROL

33 Phase 3 entry criteria

PHASE 3 AUTHORIZED

The architecture is ready to govern production of the leadership prompt library. No domain prompt may bypass DIRECT, capability truthfulness, validation gates, or the cyber control plane.

Gate	Status	Acceptance evidence
G1 — Product terminology	PASS	ChatGPT and API controls separated; High and Pro correctly represented.
G2 — DIRECT architecture	PASS	Six fields defined with templates and examples.
G3 — Prompt tiers	PASS	Executive Compact, Production High, and Audit-Grade templates locked.
G4 — Role model	PASS	One accountable role plus lenses, priorities, rights, and exclusions.
G5 — Efficiency controls	PASS	Reasoning selection, token layers, caching/context controls documented.
G6 — Capability controls	PASS	Skills, plugins, apps, tools, permissions, and availability separated.
G7 — Quality controls	PASS	Output contracts, validation gates, scoring, and release statuses defined.
G8 — Cyber controls	PASS	Authorization, routing, finding schema, and remediation gates defined.
G9 — Publishing schema	PASS	Prompt-card metadata and lifecycle statuses locked.

Phase 3 authorized workstreams

- Solution architecture and technical-decision prompt suite.
- Engineering leadership, code review, debugging, implementation, and release prompt suite.
- Product, design, frontend, accessibility, and artifact-production prompt suite.
- Data, integrations, automation, operations, and executive leadership prompt suite.
- Defensive cybersecurity prompt suite developed under the Phase 2 cyber control plane.

Change-control rule

If official OpenAI product behavior changes, update the source lock and the affected Phase 2 control before changing downstream prompts. Domain prompts inherit these controls; they do not silently redefine them.

SOURCES

A Official source register — model and prompting

S01	**OpenAI API — Model guidance: Using GPT-5.6** https://developers.openai.com/api/docs/guides/latest-model Model family, outcome-focused prompting, reasoning effort, Pro mode, verbosity, caching, persisted reasoning, programmatic tool calling, multi-agent beta, autonomy boundaries.
S02	**OpenAI API — Models catalog** https://developers.openai.com/api/docs/models Sol/Terra/Luna positioning; supported reasoning values and tool categories.
S03	**OpenAI Help Center — GPT-5.6 in ChatGPT** https://help.openai.com/en/articles/20001354-gpt-56-in-chatgpt ChatGPT Medium, High, Extra High, and Pro behavior and plan/workspace dependence.
S04	**OpenAI Help Center — Skills in ChatGPT** https://help.openai.com/en/articles/20001066-skills-in-chatgpt Skills as reusable workflows with instructions, examples, supporting resources, and code.
S07	**OpenAI Help Center — Plugins in ChatGPT and Codex** https://help.openai.com/en/articles/20001256 Plugins as packages of Skills, apps, and app templates; role/workspace dependencies.
S08	**OpenAI Help Center — Apps in ChatGPT** https://help.openai.com/en/articles/11487775 Connected data/actions, app permissions, approval controls, and source-system access limits.

SOURCE DATE

Source register verified July 26, 2026. Product availability, limits, beta status, and policy language are volatile and must be rechecked before final publication.

SOURCES

B Official source register — cybersecurity and governance

S05 — **OpenAI Deployment Safety Hub — GPT-5.6 System Card**
https://deploymentsafety.openai.com/gpt-5-6
Preparedness classification, cyber capability findings, safeguards, allowed defensive uses, and restricted harmful uses.

S06 — **OpenAI Help Center — Trusted Access for Cyber overview**
https://help.openai.com/en/articles/20001258-openai-daybreak-trusted-access-for-cyber-overview
Authorized defensive scope, eligibility, governance, and continuing policy/access requirements.

S09 — **OpenAI Help Center — Admin controls, security, and compliance for plugins and apps**
https://help.openai.com/en/articles/11509118-admin-controls-security-and-compliance-for-plugins-and-apps
Workspace roles, app access, action controls, and administrative permissions.

S10 — **OpenAI Help Center — Additional safety checks for biological and cybersecurity requests**
https://help.openai.com/en/articles/20001326-additional-safety-checks-for-biological-and-cybersecurity-requests-in-chatgpt-codex-and-the-api
Additional automated checks, possible delay/blocking, and interpretation.

S11 — **OpenAI — Cybersecurity in the Intelligence Age**
https://openai.com/index/cybersecurity-in-the-intelligence-age/
Defensive-access strategy, resilience, visibility, control, and responsible capability expansion.

S12 — **OpenAI — Terms and policies index**
https://openai.com/policies/
Current entry point for Usage Policies and governing terms.

Publication maintenance protocol

1. Recheck model guidance, ChatGPT model availability, Skills/plugins/apps documentation, system card, Trusted Access, additional safety checks, and Usage Policies.
2. Record access date and update any volatile model, plan, beta, price, permission, or policy statements.
3. Re-run representative prompt evaluations when model, prompt, Skill, tool, or source behavior changes.
4. Preserve historical versions of this control document so downstream prompt changes remain traceable.

> **PHASE 2 RESULT**
> The prompt architecture, configuration method, validation system, and defensive cybersecurity control plane are locked and ready for Phase 3 domain prompt production.